AWS
re:Invent

Agenda

- AWS Inf1 – Review from CMP324 [35 min]

- Lab Logistics and Setup: [5 min]

- Lab 1-4 [80 min]

- Learning Resources and Summary

More machine learning happens on AWS than anywhere else

Inference accounts for the majority of machine learning infrastructure costs

% infrastructure cost

Machine learning training (<10%)

Machine learning inference (>90%)

AWS Inferentia

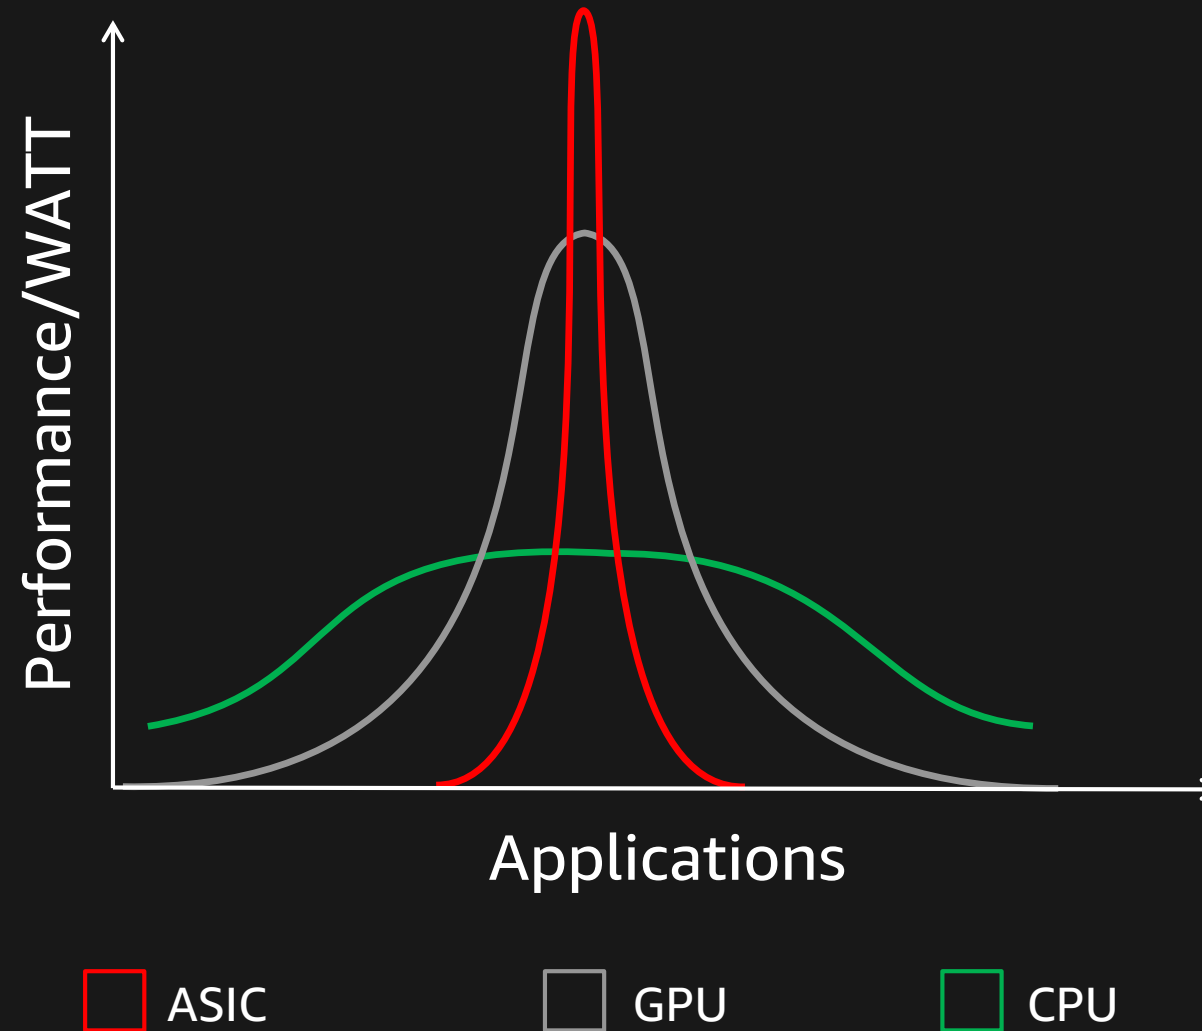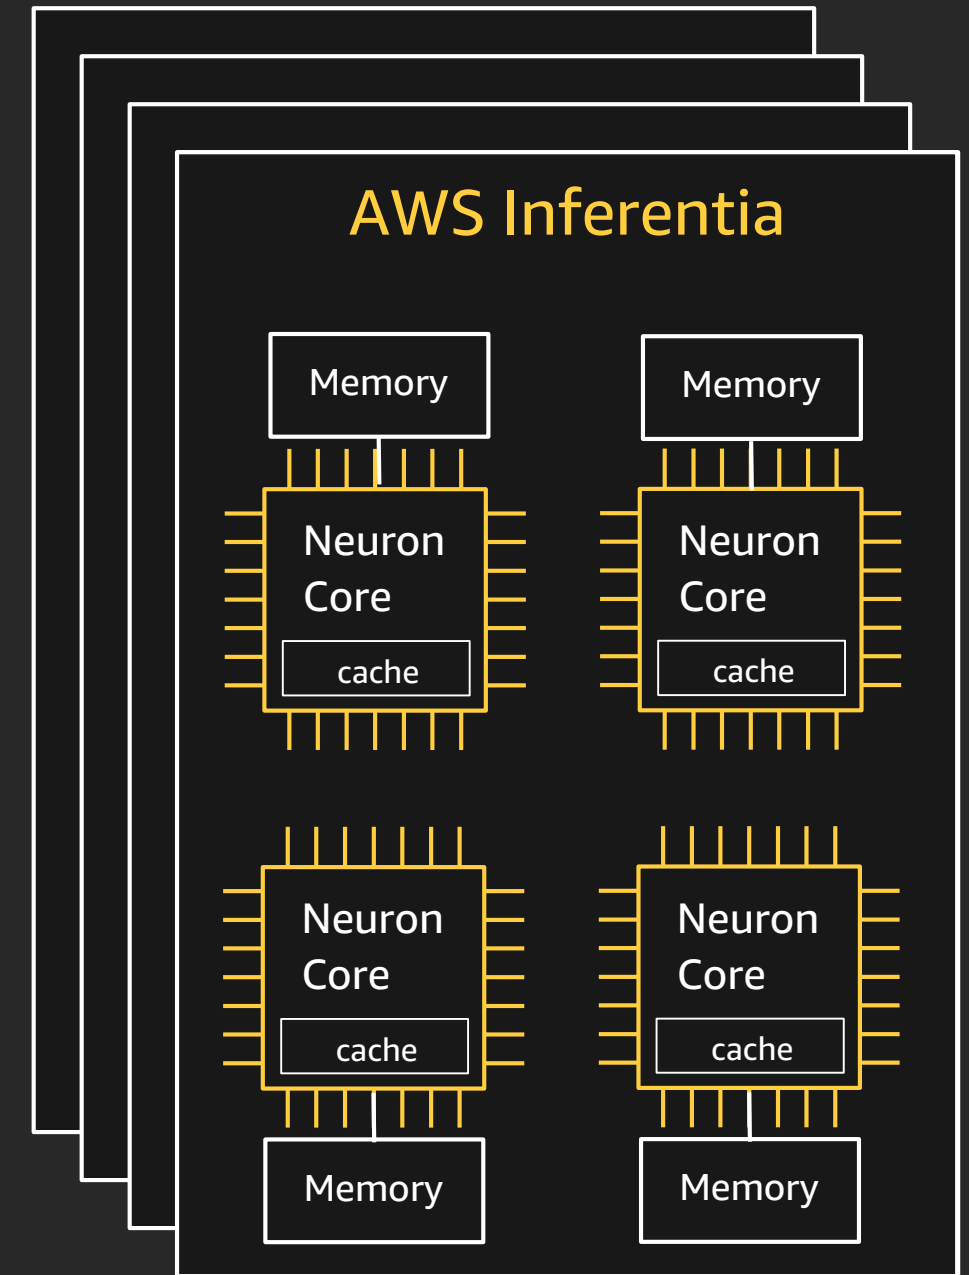# Optimizing ML performance with a custom chip


Chart for illustrative purposes

AWS Inferentia quick tour

AWS custom built: chip, software, and server

- 4 Neuron Cores

- Up to 128 TOPS

- 2-stage memory hierarchy
  - Large on-chip cache and commodity DRAM

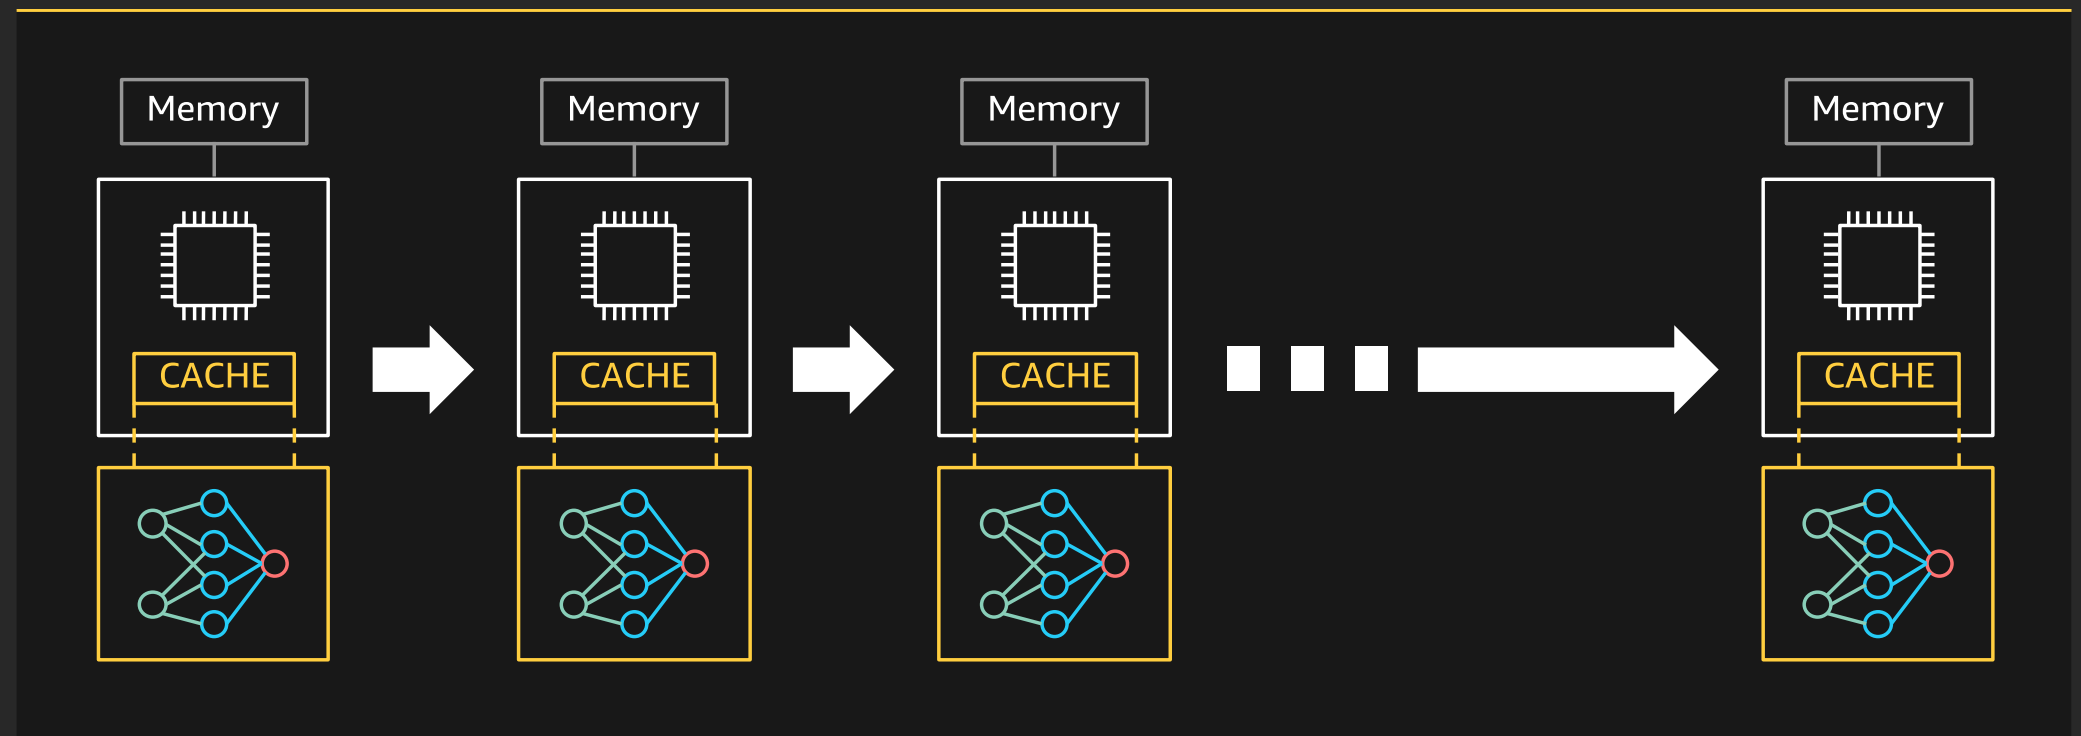- Supports FP16, BF16, INT8 data types

- Fast chip-to-chip interconnect

Low-latency inference at high load

NeuronCore pipeline

- Very low latency
- Full bandwidth due to on-chip cache

AWS Neuron

Introducing AWS Neuron

Software suite enabling high-performance deep learning inference on AWS Inferentia

Compiler

Runtime

Profiling and debugging tools

Supports all major frameworks

github.com/aws/aws-neuron-sdk

AWS Neuron support forum

AWS Neuron quick tour

Compile

Neuron Compiler
(NCC)

Deploy

Neuron Runtime
(NRT)

Neuron Binary
(NEFF)

Profile

Neuron Tools

```
C:\>code --version
1.1.1
```

AWS Neuron highlights

Smart partitioning

Automatically optimize neural-net compute

AWS Neuron highlights

Smart partitioning

Automatically optimize neural-net compute

Auto FP32 casting

Ingest FP32 trained models and Neuron auto-casts to BF16

AWS Neuron highlights

Smart partitioning

Automatically optimize neural-net compute

Auto FP32 casting

Ingest FP32 trained models and Neuron auto-casts to BF16

NeuronCore pipeline

Very low latency Full bandwidth

AWS Neuron highlights

Smart partitioning

Automatically optimize neural-net compute

Auto FP32 casting

Ingest FP32 trained models, and Neuron auto casts to BF16

NeuronCore pipeline

Very low latency Full bandwidth

NeuronCore Groups

Concurrently run multiple models
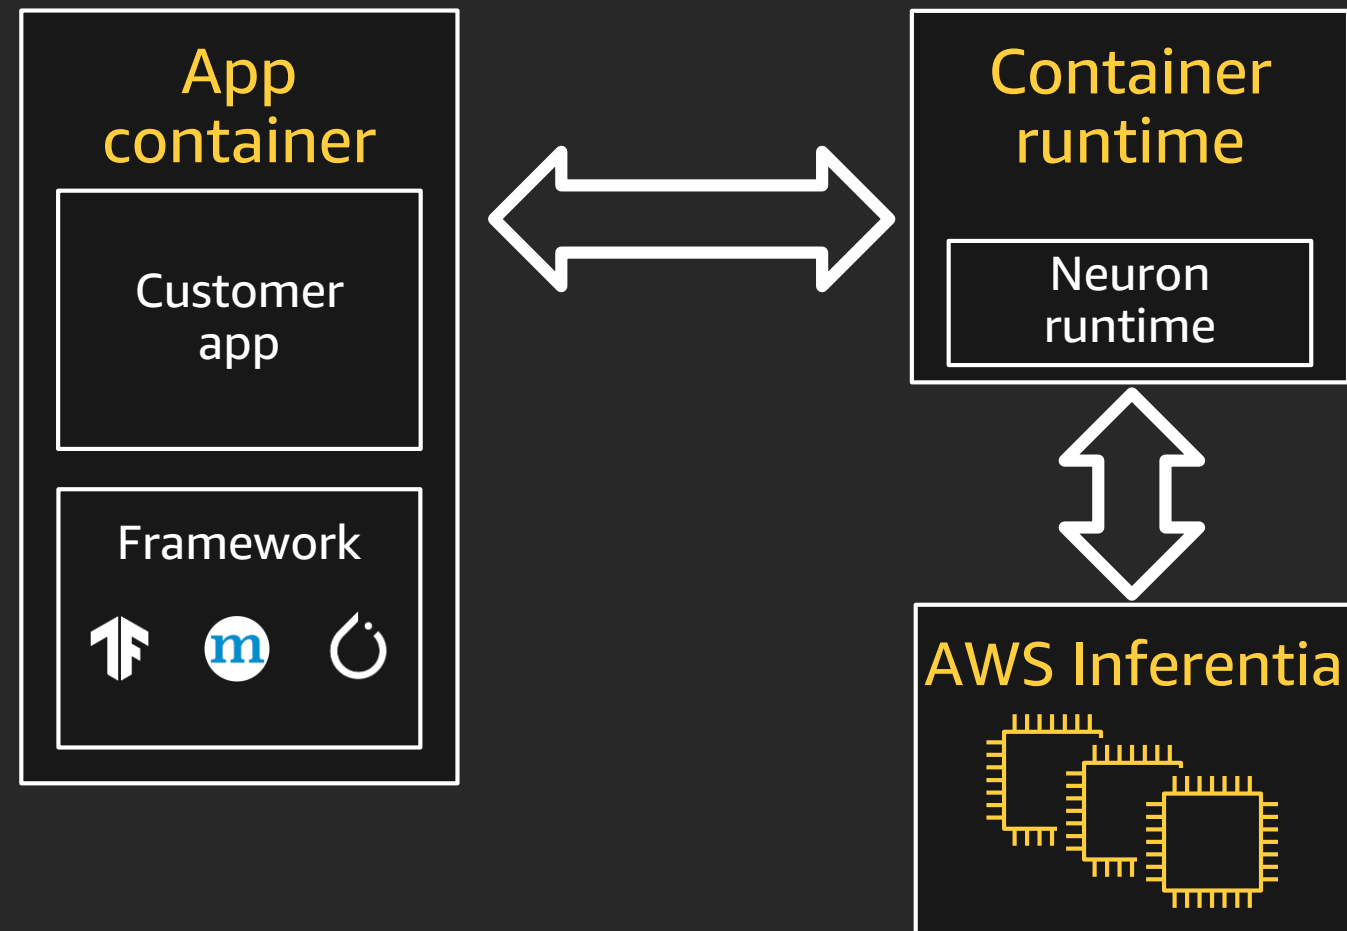
Profiling with Neuron tools

Profile
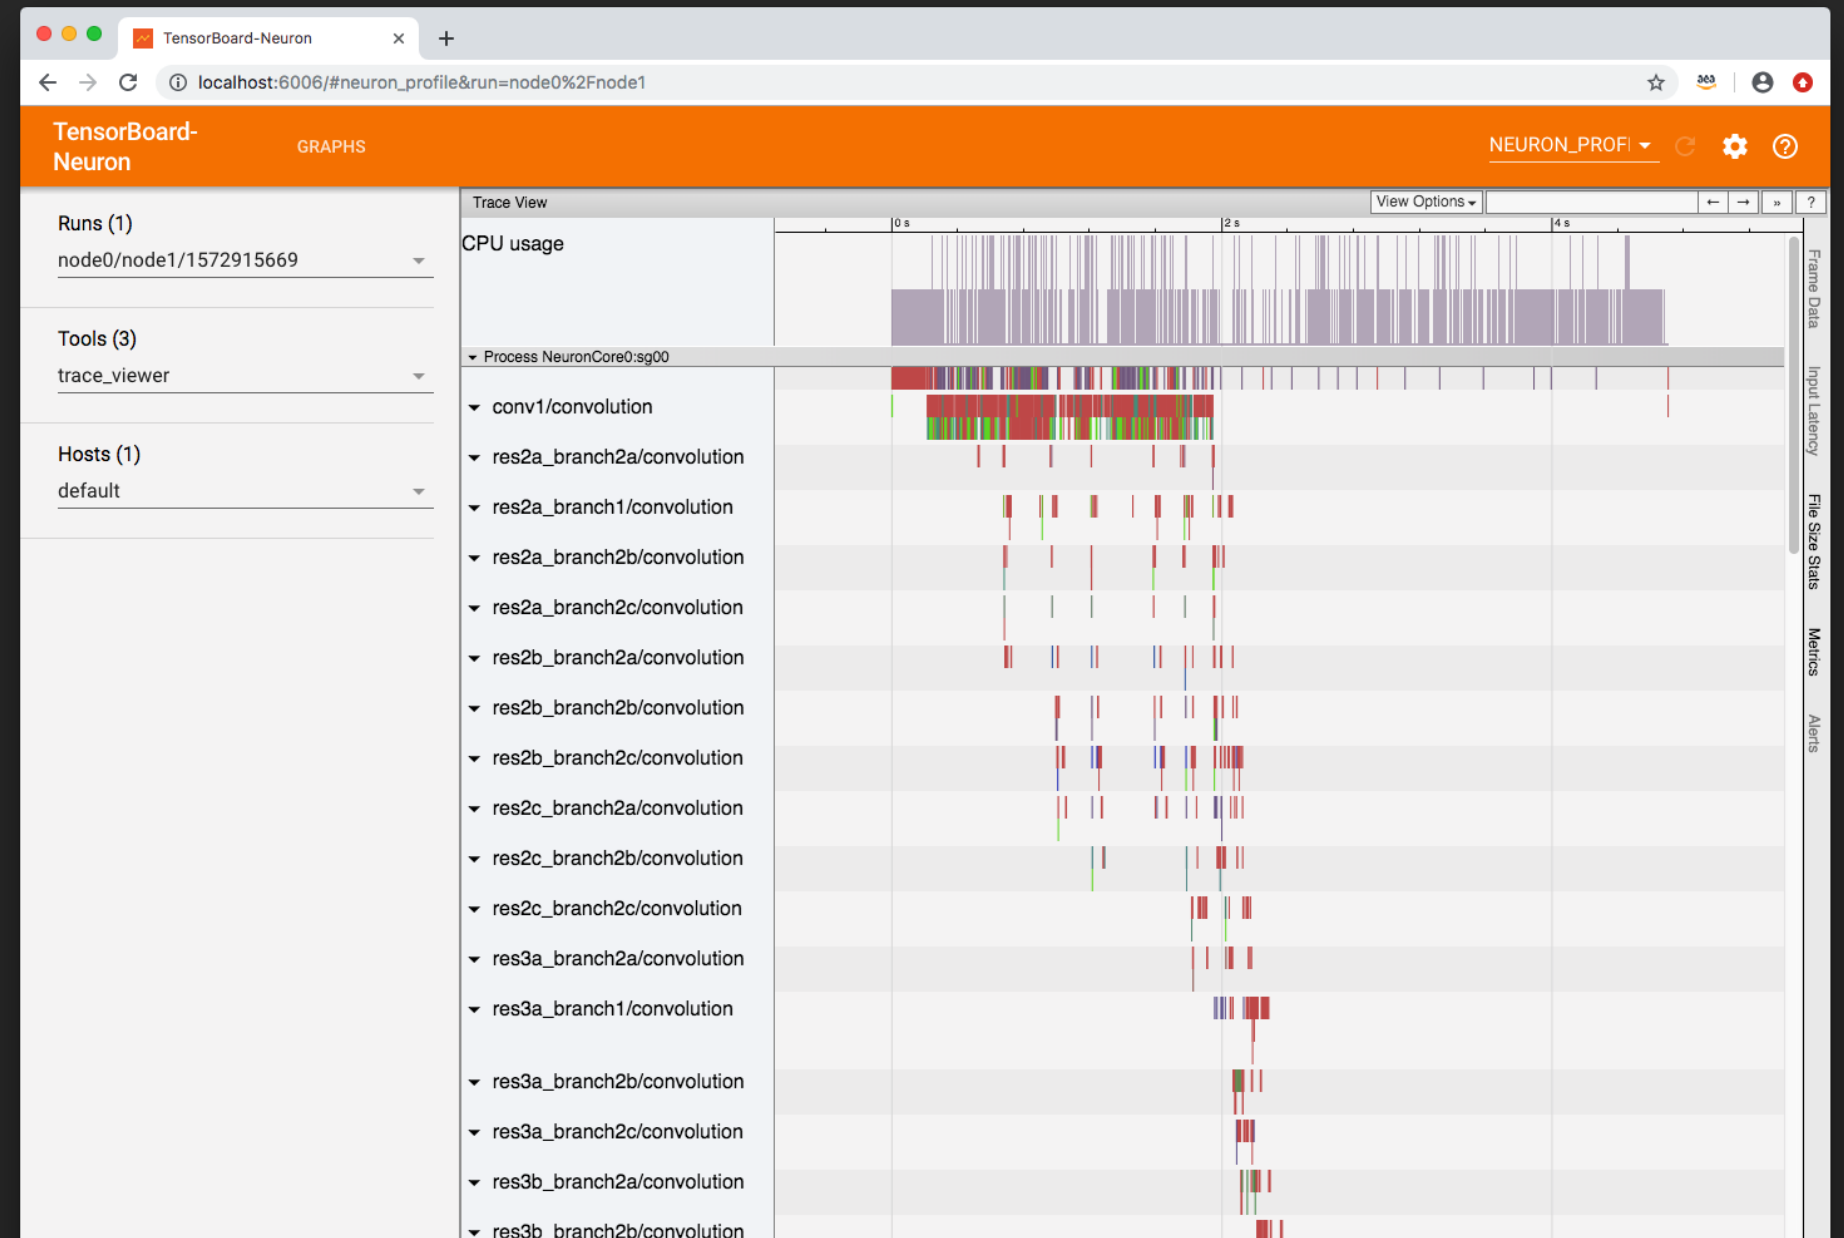
Amazon EC2 Inf1 instances

aws

Introducing Inf1 instances for ML inferencing

Object detection

Natural language processing

Personalization

Speech recognition

Image processing

Fraud detection

ML inference deployment options on Amazon EC2

Custom chip
EC2 Inf1 instances

Featuring AWS Inferentia

○ Applications that leverage common ML frameworks

○ Powered by AWS Inferentia

○ Best price/performance for ML inferencing in the cloud

Up to 40% lower inference than EC2 G4 instances

○ **Available today**

GPU based
EC2 G4 instances

○ Applications that require access to CUDA, CuDNN or TensorRT libraries

○ Amazon EC2 G4 instances based on NVIDIA T4 GPUs

○ **Launched**

CPU based
EC2 C5 instances

○ Low sensitivity to inference performance

○ Intel Skylake CPUs Support for AVX-512/ VNNI instruction set

○ **Launched**

Inf1 instance sizes

Instance size	vCPUs	Memory (GIB)	Storage	Inferentia Chips	NeuronCore Pipeline Mode	Network B/W	EBS B/W
inf1.xlarge	4	8	EBS only	1	N/A	Up to 25 Gbps	Up to 3.5 Gbps
inf1.2xlarge	8	16	EBS only	1	N/A	Up to 25 Gbps	Up to 3.5 Gbps
inf1.6xlarge	24	48	EBS only	4	Yes	25 Gbps	3.5 Gbps
inf1.24xlarge	96	192	EBS only	16	Yes	100 Gbps	14 Gbps

- Available in 4 sizes
- Single and multi-chip instances

- AWS 2nd Gen Intel Xeon Scalable Processors
- Up to 100 Gbps networking bandwidth

- Available to use with Amazon SageMaker, Amazon Elastic Kubernetes Service, and Amazon Elastic Container Service (coming soon)

Best price/performance for ML inference in the cloud

Instance type	Throughput (Seq/Sec)	OD Price ($/Hr)	Cost per inference	Throughput: Inf1 vs. G4	Cost-per-inference: Inf1 vs. G4
inf1.24xlarge	19,200	$7.619	$0.0003968		
				3.28x	40.7%
G4.12xlarge	5,846	$3.912	$0.0006692		

Results based on running BERT-base model end-to-end with TensorFlow

EC2 Inf1 offers up to 3x higher throughput and 40% lower cost per inference than EC2 G4 instances

Highest price-to-performance in the cloud for ML inference

Try Inf1 instances today

Inf1 instances are now available in the US-East-1 (N. Virginia) and US-West-2 (Oregon) Regions
More Regions coming soon!

EC2 Inf1 instances support on demand, reserved and spot purchasing options; also available as part of Savings Plan

Instance size	On Demand	1-Year Standard RI (40% discount)	3-Year Standard RI (60% discount)
inf1.xlarge	$ 0.368/Hr	$ 0.221/Hr	$ 0.147/Hr
inf1.2xlarge	$ 0.584/Hr	$ 0.351/Hr	$ 0.234/Hr
inf1.6xlarge	$ 1.905/Hr	$ 1.143/Hr	$ 0.762/Hr
inf1.24xlarge	$ 7.619/Hr	$ 4.572/Hr	$ 3.048/Hr

Alexa use case

aws

Neural text-to-speech challenges

Low latency requirement for dialog system
High throughput requirement implied by streaming of output speech

Context generation is a sequence-to-sequence auto-regressive model
Inference-bound memory bandwidth

High temporal density of speech production model resulting in 90 GFLOPs to generate 1 second of output
Compute-bound inference

Using EC2 GPU instances to meet requirements results **in high operational cost**

Alexa TTS migration to EC2 Inf1: Ease of integration

Alexa TTS uses MXNet
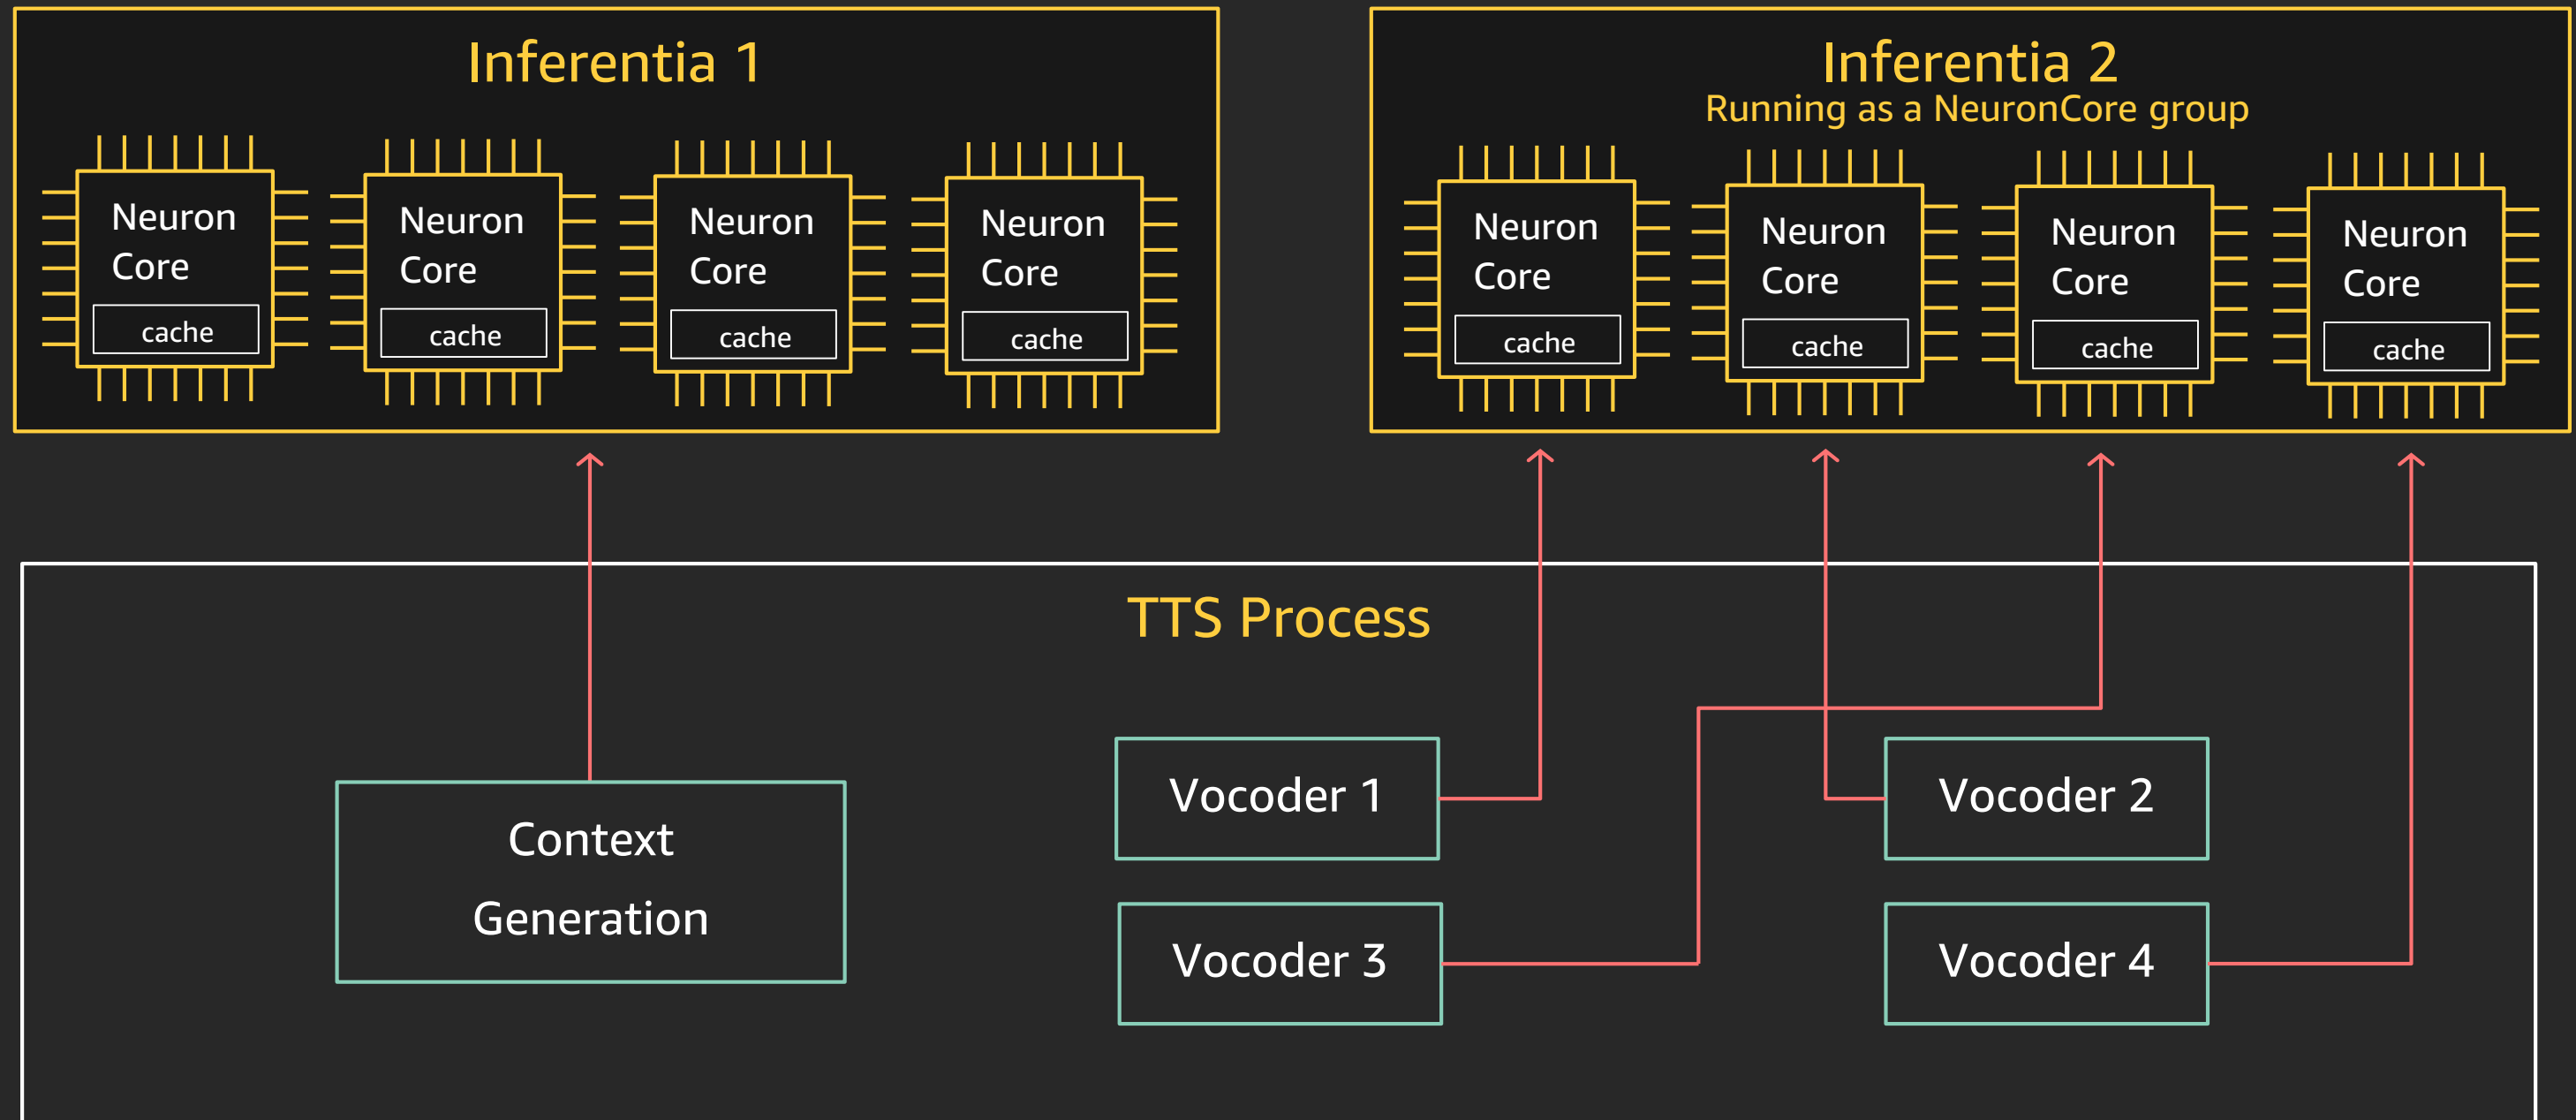supported natively by
AWS Neuron

Support for
C and Python
APIs

Options to migrate
original FP32 models
to FP16 or Bfloat16

Alexa TTS migration to EC2 Inf1: Architecture

Alexa TTS migration to EC2 Inf1: Long text

**Alexa long text traffic (ex: books, news) has even higher gains**

120%

100% — **100%**

80%

60% — Lowest cost

40% — **45%**

20%

0%

Speech Generation Cost

74%

■ P3   ■ G4   ■ Inf1

Getting started

Github
Tutorials/drivers

**AWS Neuron
developer forum**

**EC2 Inf1
landing page**

Hands-on labs

1. Set up development instance on a **c5d.4xlarge**

2. Deploy and model-serve on Inf1.2xlarge instance

3. Load test run

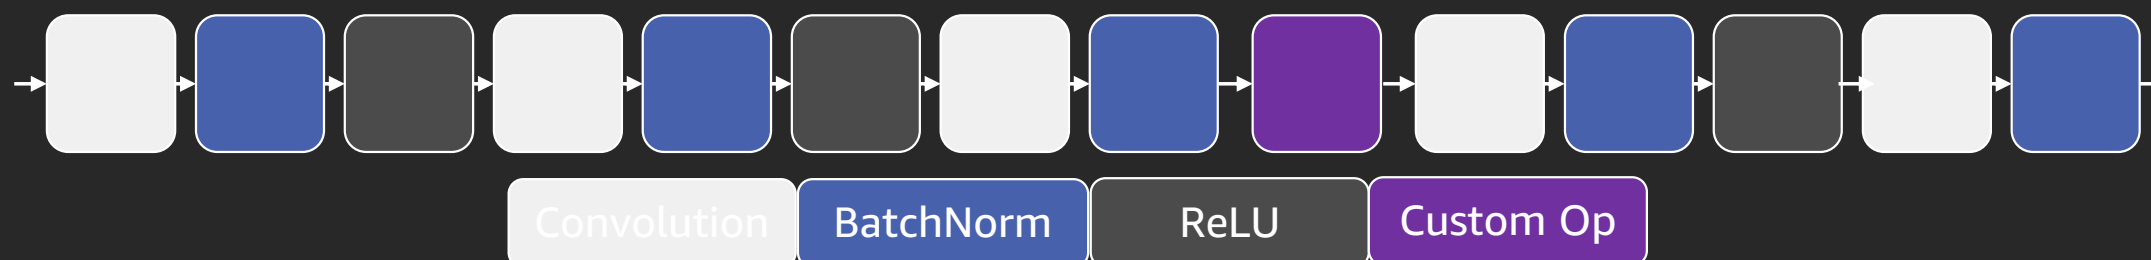4. Debugging and profiling

Full URL:

https://github.com/awshlabs/reinvent19Inf1Lab

Questions? Use Forum at:
https://forums.aws.amazon.com/forum.jspa?forumID=355

Thank you!

Subgraph Compilation – under the hood

- We start with a normal computation graph for an MXNet model

| Convolution | BatchNorm | ReLU | Custom Op |

- We search for subgraphs of supported operators

| Supported Op | Custom Op |